

Lifecycle Management For Docker UI

FINAL REPORT (SEMESTER 2)

Stephen Coady

20064122

Supervisor: Dr. Brenda MULLALLY

BSc (Hons) in Applied Computing

Plagiarism Declaration

Unless otherwise stated all the work contained in this report is my own. I also state that I have not submitted this work for any other course of study leading to an academic award.

Contents

1	Introduction	5
1.1	Problem Statement	5
1.2	Aims and Objectives	6
2	Technologies	8
2.1	Docker	8
2.2	Node JS	8
2.3	Angular JS	9
2.4	Travis CI	9
2.5	Vagrant	10
3	Design	11
3.1	System Architecture Overview	11
3.2	Front End Design	11
3.3	Enterprise Level Concerns	11
4	Methodology	12
4.1	Agile	12
4.2	Jira	12
4.3	Gitflow	12
4.4	Testing	12
4.5	Continuous Integration/Deployment	12
	Appendices	13
A	Application Repository	13
B	Dockerode	13
C	Travis Repository	13
D	SonarQube Repository	13
E	Sprint Retrospectives	13
	Bibliography	17

Glossary

- API** Application programming interface. A set of endpoints which lead to functions and application logic. Allows developers to expose application functionality without directly exposing the code within the application. [5](#)
- CLI** Command Line Interface. A means of interacting with a computer program where a user can only enter commands in the form of successive lines of text. [5, 6](#)
- continuous integration** The act of continuously merging newly created software with the current working version, validated by automatic build and testing tools allowing for early detection of faults. [9](#)
- Docker** An application which manages containers running on a host. [5, 8](#)
- Docker container** An isolated package which bundles everything required to run an application within a container. The application can then be run within this container without regard to the underlying architecture. [8](#)
- Docker daemon** The server-side component of the Docker Engine. [6](#)
- Docker host** The computer, server or virtual machine on which the Docker application is installed. One user may manage several Docker host's, depending upon how many computers they have installed it on. [6, 8](#)
- Docker image** A template from which many containers can be started. Can be pushed and pulled to/from a remote Docker registry. [6, 8](#)
- REST** Representational state transfer. A means of providing communication between computers across the internet. RESTful web services allow systems interact with each other using completely stateless operations. [6](#)
- Travis** A website which allows for automatic builds of code and feedback of results. Can be used to automatically deploy code dependent on test results also. [9](#)
- UI** User Interface. A graphical view towards an application which exposes functionality by using buttons or other components. Typically utilised using a computer mouse. [5](#)
- Vagrantfile** A template file which provides Vagrant with all the instructions required to create the desired environment. Written in the Ruby programming language. [10](#)

virtual machine An abstracted virtual version of a computer. Normally runs on another host with a dedicated piece of software to manage it. Appears to the user as a real computer, however all or most components are defined by software. [8](#), [10](#)

1 Introduction

This report will aim to guide the reader through the planning and development of the Lifecycle Management for Docker UI application. After reading this report the reader should have a clear idea of why the application was built, what was used to build it and how the process was carried out.

This application will be built using open source principles and best practices, enabling it to be maintained and improved by any developer who wishes to contribute. For this reason many of the decisions made and processes employed were done so with an open source final product in mind.

1.1 Problem Statement

Currently the Docker application does not ship with any bundled UI. When installed, it is comprised of a client and a server side component (Docker, 2017). The server side exposes itself through an API and is ultimately responsible for controlling all aspects of Docker on the host such as containers, images, networks and volumes etc. The API exposed by the server-side application of the Docker Engine is consumed by the Docker CLI which is the client side application. A graphical representation of the complete Docker engine can be seen in Figure 1.

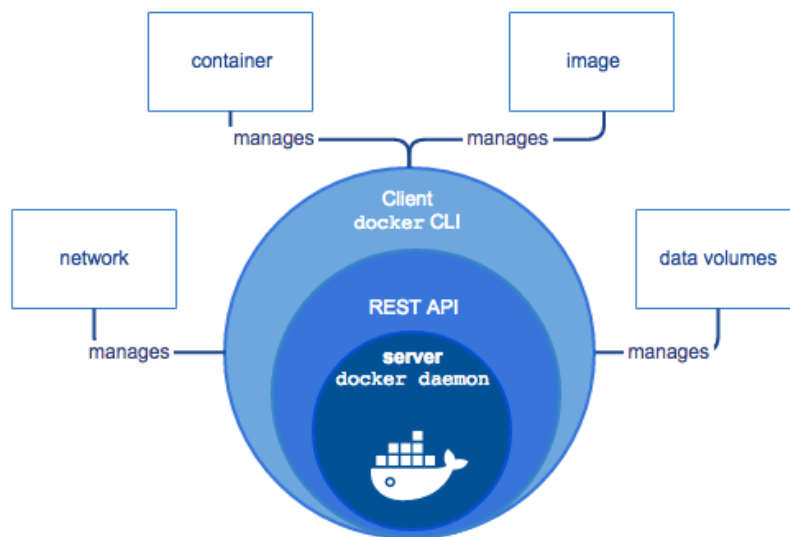


Figure 1: *Docker Engine Components.* Credit: (Docker, 2017)

This model is extremely versatile as it allows the developer to control any [Docker daemon](#) (the server-side component of a Docker installation) once they have access to the command line of the host Docker is running on. In fact, if the API exposed by the Docker daemon is exposed remotely then the developer does not need access to the host's command line, instead they can directly access the API remotely.

While the [CLI](#) gives developers full control over the server-side component of a [Docker host](#) it also has its drawbacks.

- Learning curve - the person using the command line must be familiar with typical commands used to achieve certain tasks. This precludes anybody without these skills from using Docker.
- Vast set of Docker commands - there are a vast number of commands available to use from the Docker CLI. This is also a learning curve as even a developer who is familiar with a CLI must first learn the Docker commands to be able to use the client-side application.
- User friendliness - The command line does not produce content that is easily readable and can often format the data it is trying to present in an odd fashion depending on things like screen size etc.

1.2 Aims and Objectives

The aim of this project is to address all of these problems while also trying to increase the functionality available to anybody who wishes to use Docker.

At a high level the primary objectives of this project are:

- Fully functional server-side application
- Expose this application through a [RESTful](#) API.
- A fully functional front-end application
- A [Docker image](#) built to allow easy distribution of the application

These objectives will then provide the following functionality:

- A UI which will
 - allow users to manipulate images, containers etc on the host with the same capabilities as the command line
 - allow them to do this remotely
- A runnable container which has no other dependencies so that it can be run without installing the application

- An independent API which can be consumed by any front end application
 - This will provide flexibility if the front end framework needs to be changed further down the line

2 Technologies

In this section of the report the technologies used to create the application will be examined and explained in detail.

2.1 Docker

[Docker](#) is a platform which allows developers to package their applications into isolated containers which contain only the software dependencies required by the application. A container is different to a [virtual machine](#) in that a container does not contain a full operating system ([Docker, 2016](#)).

Since a primary objective of this application is to manage a [Docker host](#) it makes sense to leverage the capabilities of Docker and run this application within a [Docker container](#). This provides several benefits over distributing source code, such as:

- Portability - If a [Docker image](#) can be built and uploaded to a public repository then it makes it easier for other developers to pull and run the application.
- Ease of use - As the application will be running in a container a developer does not need to install any third party components on their system to use the application. They do not need to worry about their environment at all, once their system has Docker installed it will run the application.
- Ephemeral - Docker containers are designed to be ‘throw-away’. This means that if this application needs to be quickly stopped and restarted then containers are the perfect vehicle to do this.

2.2 Node JS

Node JS is a server-side JavaScript runtime, it is built on the same V8 engine that powers the popular Chrome browser. It uses an event-driven, non-blocking I/O model that makes it lightweight, efficient and very fast. Node JS’ package ecosystem, NPM, is the largest ecosystem of open source libraries in the world. ([Nodejs.org, 2016](#)). Node JS provides an excellent way to build highly scalable network application which are non blocking and extremely performant ([Griffin et al., 2011](#)). This means that if the application was adapted in the future to deal with large numbers of servers then the technology choice will be able to deal with that.

Node JS was also deemed a good fit for this project as it has a large and extremely active online community. Since this is an open source project this will increase the likelihood of other developers taking part in the project and contributing. We can see in [Figure 2](#) that there are currently (as of

November 2016) more node modules available through the node package manager (NPM) than any other of the large package managers such as the ones used by Go, PHP, Python and Ruby.

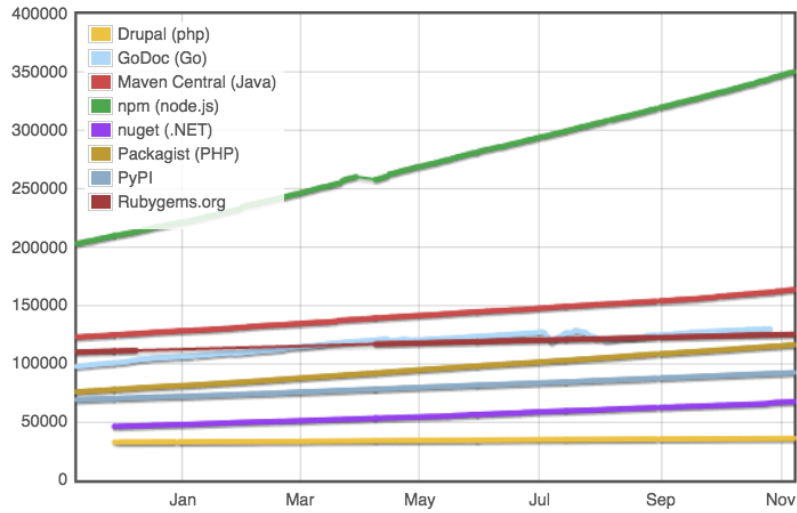


Figure 2: *Various Module Counts (Credit: modulecounts.com)*

2.3 Angular JS

Angular JS is a web framework for building dynamic web applications using Javascript as a controller. It uses HTML as its template language to display the information passed to it from the application controller ([AngularJs, 2017](#)).

For this application the front end requirements were relatively low. Once the framework provided a mechanism to show dynamic content then it was a candidate. For this reason several were evaluated and Angular was chosen as it had the lowest barrier to entry. The front-end application will be discussed further in Section 3.

2.4 Travis CI

Since this project was developed as an open source project it made sense to have a transparent and fully automated build process integrated into the project. For this reason [continuous integration](#) in this project is handled by [Travis CI](#). It is a feature-rich service free to use for open source projects and has a large user base. The set up used in this project for testing and continuous integration will be explored further in Section 4.

2.5 Vagrant

Vagrant is a technology to create configurable, reproducible and portable environments by using a set of programmable steps to produce a [virtual machine](#) which the developer can use to develop in ([Vagrant, 2017](#)). While this is just one use-case of Vagrant it is the main reason it was used in this project.

Since this project is open source it is useful to have one standardised VM within which all development can take place. This virtual machine can then be shared (either as its own separate repository or included in the main application repository). This is useful as it enables any developer who wishes to contribute to the project to instantly have the required environment. For instance, a developer can contribute to this project by using the [Vagrantfile](#) supplied without requiring them to first install Docker or Node JS.

3 Design

3.1 System Architecture Overview

3.2 Front End Design

3.2.1 Wireframes

3.2.2 Mobile Web Application

3.3 Enterprise Level Concerns

Performance:

- load balancing
- scaling
- caching

4 Methodology

4.1 Agile

4.2 Jira

4.3 Gitflow

4.4 Testing

4.5 Continuous Integration/Deployment

Appendices

A Application Repository

<https://github.com/StephenCoady/lifecycle-management-for-docker>

B Dockerode

<https://github.com/apocas/dockerode>

C Travis Repository

<https://travis-ci.org/StephenCoady/lifecycle-management-for-docker>

D SonarQube Repository

<https://sonarqube.com/dashboard/index?id=lifecycle-management-for-docker>

E Sprint Retrospectives

2017-02-01 Docker Project Sprint 1 Retrospective

Date 01 Feb 2017

Participants [Leigh Griffin](#) [Stephen Coady](#)

Retrospective

What did we do well?

- Already had prototype in place to accelerate development
- 3rd party module knowledge accelerated development
- Guidance from Red Hat really helped focus the sprint
- Scope on tickets well understood from Red Hats perspective

What should we have done better?

- Story points were inflated because domain knowledge was higher than anticipated
- Testing strategy needs to be revised, very time consuming

Actions

- [Stephen Coady](#) review backlog for story point accuracy based off of current domain knowledge
- [Stephen Coady](#) add a ticket to review / spike testing strategies, feel free to consult [David Martin](#) and [Leigh Griffin](#) on specifics
- [Stephen Coady](#) add a ticket for UI frameworks investigations and spikes, end result should be an Epic that we can triage and prioritise

2017-02-15 Docker Sprint 2 Retrospective

Date 15 Feb 2017

Participants [Leigh Griffin](#) [Stephen Coady](#) [David Martin](#)

Retrospective

What did we do well?

- We performed a backlog review and set the priority for the remaining 3 sprints
- Progression of the sprint was excellent, good pace to it and good story pointing
- Comms was pretty good
- Smooth sprint, story points and tickets were well scoped
- Sprint Planning revisited the story points so very little surprises
- Team (Stephen) came to the Stakeholders (Dave & Leigh) with the plan for the next Sprint

What should we have done better?

- Sprint started at a bad time college wise, with other assignments due
- Not keeping in touch with the sprint day to day (Dave & Leigh)

Actions

- [Stephen Coady](#) to share wireframes as a mid sprint review asynchronously. Would recommend emailing that to us both.
- [Stephen Coady](#) metrics spike insight when you get to it (this sprint possibly)

2017-02-28 Docker Sprint 3 Retrospective

Date 28 Feb 2017

Participants [Leigh Griffin](#) [Stephen Coady](#)

Retrospective

What did we do well?

- The UI is very usable, lots of nice feedback and functionality visible now
- Consistency in the velocity at 20 points
- Got the wireframe relationship, it really helped with my front end skills which I was not confident in
- Wireframe feedback was excellent, really helped scope the work
- Got to demo to my supervisor which gave her a lot of insight
- Overall work pace was judged well for the most part, consistent delivery

What should we have done better?

- The story pointing on the skeleton was completely off, it could have derailed the entire sprint
- Velocity last sprint was off
 - you should have descoped when you realised how big the UI was
 - you should have re story pointed the UI mid sprint to allow a controlled descope
- Tickets are not descriptive enough, need to add more metadata
- Didn't descope the testing in a container ticket, should have done that when the UI became so big

Actions

- [Stephen Coady](#) to review the backlog with a view to WHAT and WHY being evolved in the tickets as well as story points
- [Stephen Coady](#) to define the critical path through the project, ~80 story points left with a ~60 story point burn predicted
- [Stephen Coady](#) to add some investigative tasks around KeyCloak SSO for future work i.e. out of scope of this project

Bibliography

AngularJs (2017), ‘What is angularjs?’. [online] Accessed: 10-03-2017.

URL: <https://docs.angularjs.org/guide/introduction>

Docker (2016), ‘What is Docker’. [online] Accessed: 27/11/2016.

URL: <https://www.docker.com/what-docker>

Docker (2017), ‘Docker overview’. [online] Accessed: 08/03/2017].

URL: <https://docs.docker.com/engine/understanding-docker/>

Griffin, L., Ryan, K., de Leastar, E. and Botvich, D. (2011), ‘Scaling instant messaging communication services: A comparison of blocking and non-blocking techniques’. [online] Accessed: 10-03-2017.

URL: <http://repository.wit.ie/1636/>

Nodejs.org (2016), ‘Node.js’. [online] Accessed: 27/11/2016.

URL: <https://nodejs.org/en/>

Vagrant (2017), ‘Why vagrant?’. [online] Accessed: 10-03-2017.

URL: <https://www.vagrantup.com/docs/why-vagrant/>